

Neural Inference as the Operating System on Minimal RISC-V

Bernhard Guillon

Advanced Systems Engineering (ASE2026)

Paris Lodron University of Salzburg

Salzburg, Austria

Email: Bernhard.Guillon@stud.plus.ac.at

Abstract—This paper presents an end-to-end neural execution stack on a minimal RISC-V environment, where neural networks serve as the primary computation model instead of traditional machine code. We describe a complete toolchain from PyTorch model training through custom assembly to execution on both a C++ emulator and Verilator RTL backend. The system introduces descriptor-based neural custom instructions for MLP inference, block-diagonal model composition via a glue JSON format and Rust compiler, and OS-like task switching through a learned router MLP. Evaluation demonstrates up to 2000x speedup over software-only baselines while maintaining deterministic output across dual backends through automated differential validation.

Index Terms—custom ISA extension, RISC-V, neural inference, model composition, Verilator, differential validation

PROJECT CONTEXT

Supervisor: Univ.-Prof. Dipl.-Inform. Dr.-Ing. Christoph Kirsch

Course: ASE2026 (Advanced Systems Engineering)

University: Paris Lodron University of Salzburg

Date: 2026-06-18

A. Introduction

Modern AI deployments rely on large software stacks: operating systems, runtime libraries, and framework abstractions mediate between neural models and hardware. This paper investigates the opposite direction: how far can a compact execution stack be pushed when neural operations are represented as first-class ISA extensions on a minimal RISC-V core?

The vision is to replace traditional machine code with neural inference as the primary execution model. A 141-line `runtime.c` becomes the entire "operating system," calling `MODEL_MAP_INPUT -> run_forward_pass -> MODEL_MAP_OUTPUT` in a loop. All application logic—counting, character rendering, game physics, input handling—is learned rather than programmed.

We make four contributions:

1. **Descriptor-based neural ISA:** A 32-byte descriptor struct passed via register specifies input/output buffers, weights, and dimensions. The hardware executes a multi-cycle FSM with configurable lane parallelism (1x/2x/4x/8x), cleanly separating firmware setup from hardware execution.

2. **Block-diagonal model composition:** Independently trained MLPs are merged into a single weight matrix via a declarative glue JSON format. A Rust compiler reads the glue, assembles block-diagonal weights, and emits C macros for input/output mapping. No retraining is required.
3. **Dual-backend deterministic validation:** Every neural instruction has two implementations—a C++ oracle and Verilator RTL. Automated CI tests enforce bit-exact byte-level parity between them.
4. **OS-like neural task switching:** A learned router MLP switches between sub-models based on keyboard input, mimicking process scheduling. Background tasks can run with or without updates, analogous to context switching and preemption.

B. Related Work

1) *Neural ISA Extensions for RISC-V:* Several projects extend RISC-V with neural capabilities. **RV-SCNN** (IEEE TCAD 2025) uses SIMD custom ops for CNN/SNN inference on CV32E40P. **MARVEL** (arXiv 2025) generates model-class-specific extensions via ASIP Designer, achieving 2x speedup and 2x energy reduction for CNNs. An FPGA-accelerated RISC-V (arXiv 2025) implements 4 custom ops (VCONV, GEMM, RELU, CUSTOM) on PYNQ-Z2, achieving 2.14x latency reduction and 49.1% energy savings. **Vmx-dotp** (arXiv 2026) provides RVV-based Microscaling (MX) format acceleration for transformer workloads, achieving 125 MXFP8-GFLOPS. **Mixed-precision RISC-V** (arXiv 2024) introduces ISA extensions for multi-pumped soft SIMD operations, achieving 15x energy reduction for DNNs.

Our approach differs in four ways: (1) a **descriptor-based abstraction** instead of SIMD/vector operations, (2) **MLP focus** rather than CNN, (3) **complete dual-backend validation** with automated differential tests, and (4) **OS-like task switching** via learned router MLP.

2) *AI-as-OS and Bare-Metal Neural Computing:* **nCPU** (Price, 2025) implements a fully differentiable CPU where every ALU operation is a trained neural network. **embodiOS** (2025) runs an LLM as an OS on x86_64 UEFI. **OSymbiote** (2026) places an AI agent as PID1 on Linux. **OO** (Operating Organism, 2025) performs bare-metal Mamba/LLaMA inference via UEFI.

These projects share the vision of neural-first computing but differ in scale and approach. nCPU runs on GPU via Apple Silicon Metal; we target bare-metal RISC-V. embodyOS and OO use LLMs; we use lightweight MLPs with custom instructions. Our contribution is demonstrating the pattern concretely with working neural models that replace specific application functions.

3) *Model Merging and Composition*: **FS-Merge** (Shwartz-Ziv et al., 2024) produces block-diagonal merged weights as an intermediate form via learnable foldable supernets. **LoRA-LEGO** (Wang et al., 2024) proves concatenation-summation equivalence for LoRA matrices. **mergekit** (Arcee AI, 2023) provides practical tools for LLM merging including passthrough methods.

Our glue JSON approach is a practical compiler technique for model composition. Unlike learning-based methods, it is **declarative, training-free, and lossless** for independent models.

C. System Architecture

The system is a co-designed stack with six layers, from trained model to executable binary:




Fig. 1. End-to-End Neural Execution Pipeline. The pipeline ensures bit-exact parity between software and hardware backends through automated differential tests.

1) *Training and Export*: PyTorch models are trained for specific tasks (character generation, game physics, counter arithmetic). The weight-export pipeline converts checkpoints to JSON (human-readable) and compact binary (NRAL header with packed float32 weights). Weight transposition matches row-major inference access patterns.

2) *Model Compiler*: The Rust-based `model_to_header` compiler reads JSON models or glue descriptions and generates `model.h` containing:

- Model dimensions (`MODEL_INPUT_SIZE`, `MODEL_OUTPUT_SIZE`)
- State variables with initial values
- `MODEL_MAP_INPUT(buf)` macro for one-hot encoding
- `MODEL_MAP_OUTPUT(buf, fb)` macro for output decoding

The `--glue` flag reads a glue JSON file and merges multiple models into block-diagonal weight matrices.

Instruction	Opcode	Function
nmatvec	0x77	Matrix-vector multiply (baseline)
nrelu	0x77	ReLU activation
nvsigpwl	0x77	Sigmoid piecewise-linear
nvclumpu8	0x77	Clamp to uint8
nmatvec8xp4	0x7b	8-lane PMAC4 (optimized)

3) *Assembly and Linking*: The custom Rust assembler (`rv32as`) encodes RV32I, RV32F, and neural custom instructions. It provides deterministic encoding with GNU-compatible output. The CMake build system generates ELF binaries by linking `runtime.c` with the generated `model.h` and assembled model blobs.

4) *Runtime*: The 141-line `runtime.c` implements the inference loop:

```

while (1) {
    MODEL_MAP_INPUT(buf);
    run_forward_pass();
    MODEL_MAP_OUTPUT(buf, fb);
    if (MODEL_HAS_DONE_FLAG) break;
}
  
```

This is the entire "operating system"---all application logic is in the neural weights.

5) *Dual Backends*: Two execution backends provide validation:

- **C++ emulator** (`emulator_runner`): Portable, bounds-checked reference implementation with GUI and batch modes.
- **Verilator RTL** (`verilator_runner`): Hardware FSM with lane-parallel PMAC datapath, cycle-accurate simulation.

6) *Verification*: Automated CI tests enforce correctness:

- Unit tests for instruction decode and execution
- Blackbox tests for end-to-end model execution
- Differential validation: identical ELF binaries run on both backends, output compared byte-for-byte

D. Neural ISA Extension

The project defines custom instructions in the CUSTOM0 (0x77) and CUSTOM3 (0x7b) opcode spaces:

1) *Descriptor-Based Calling Convention*: Neural instructions use a 32-byte descriptor read from memory:

```

neural_desc_t desc = {
    .input_ptr = INPUT_BUF,
    .weights_ptr = weights_addr,
    .bias_ptr = biases_addr,
    .output_ptr = output_addr,
    .input_len = input_size,
    .output_len = output_size,
    .flags = 0, .reserved = 0
};
neural_matvec(&desc); // single custom instruction
  
```

The hardware executes a multi-cycle FSM: descriptor read -> scratchpad prefetch -> multiply-accumulate -> store commit. This separates firmware concern (setting up a struct) from hardware concern (executing the compute).

Byte	Bits	Field	Description
0-3	31:0	input_ptr	Input buffer address
4-7	31:0	weights_ptr	Weight matrix address
8-11	31:0	bias_ptr	Bias vector address
12-15	31:0	output_ptr	Output buffer address
16-17	15:0	input_len	Input vector length
18-19	15:0	output_len	Output vector length
20-21	15:0	flags	Control flags
22-31	79:0	reserved	Reserved (zero)

Fig. 2. Neural ISA Descriptor Layout (32 bytes). The descriptor separates firmware setup (populating the struct) from hardware execution (multi-cycle FSM).

Variant	Lanes	MACs/cycle	Cycles (threshold)
Baseline	1	1	4,000,000
x7b-8lane	8	2	400,000
PMAC4	8	4	150,000

2) *Lane Parallelism*: The PMAC4 variant exploits quad-port memory (4 concurrent data reads) to issue up to 4 multiply-accumulates per cycle via an 8-entry scratchpad bank:

The 26.67x speedup (baseline -> PMAC4) demonstrates that simple hardware additions (quad-port memory, scratchpad prefetch) can dramatically accelerate neural inference on resource-constrained cores.

E. Block-Diagonal Model Composition

1) *Concept*: Two independent MLPs can be combined into a single network by stacking their weight matrices block-diagonally:

$$W_{combined} = \begin{bmatrix} W_A & 0 \\ 0 & W_B \end{bmatrix}$$

$$W_{combined} \cdot [x_A; x_B] = [W_A \cdot x_A; W_B \cdot x_B]$$

This is mathematically exact: no approximation, no training, no alignment needed. The models run independently within a single inference pass.

2) *Glue JSON Format*: A glue file declares how to merge sub-networks:

```
"input_mapping": "combined_counter_chargen", "initial_state": {"model_counter": 97}, "models": [{"name": "count"
```

Fig. 3. Glue JSON format for merging sub-networks.

Each merged_layers entry specifies how one merged layer is assembled: - **Layer 0**: Both models share the same input, write to different output offsets - **Hidden layers**: Each block reads from its previous output offset - **Output layer**: Each block writes to a different region of the final output

3) *Rust Compiler Integration*: The model_to_header --glue command:

1. Reads the glue JSON and loads all referenced model JSONs
2. Builds merged weight matrices with block-diagonal structure
3. Emits model.h with MODEL_MAP_INPUT and MODEL_MAP_OUTPUT macros

Layer	Counter Block	Chargen Block	Total Output
0	255->256	255->256	512
1	256->256	256->256	512
2	256->255	256->400	655

Component	Neural Equivalent	OS Equivalent
Router MLP	Learned gating function	Process scheduler
Sub-models (chargen, squash)	Independent tasks	User processes
Router gates	Sigmoid outputs [0,1]	Process priority/selection
Framebuffer selection	Gate comparison	Display output routing

4. Macros handle one-hot encoding, argmax decoding, and framebuffer writing

4) *Counter+Chargen Example*: The counter+chargen combined model merges:

- **Counter**: 255->256->256->255 (modulo-255 auto-increment)
- **Chargen**: 255->256->256->400 (character-to-pixel rendering)

Into a single 3-layer network with block-diagonal weights:

Only the chargen output (400 values) goes to the framebuffer; the counter state is stored in debug memory.

F. OS-Like Neural Task Switching

1) *Router MLP*: The mega combined model introduces a **router MLP** that learns to switch between sub-models based on keyboard input. This mimics an operating system scheduler:

The router is a small 3-layer MLP: 2 -> 4 -> 2 -> 2 (relu, sigmoid, none). It is trained with identity mapping: input [1, 0] -> output [~1, ~0] (chargen active), input [0, 1] -> output [~0, ~1] (squash active).

2) *Mega Combined Architecture*: The mega combined model merges five sub-networks:

Output buffer layout (1009 neurons):

3) *Task Switching Behavior*:

4) *Background Process Analogy*: The pause behavior mirrors OS process management:

- **Escape (pause)**: Analogous to **suspending a process**. The squash game's state variables retain their last values. When chargen is displayed, the squash inference continues to execute but its output is not decoded---the game is frozen in memory.
- **Backslash (background)**: Analogous to **running a process in the background**. The squash game keeps updating its state (ball moves, paddle responds) while chargen is displayed. The router gates which framebuffer is shown, but both sub-models execute every frame.

This demonstrates that neural inference can implement process scheduling semantics without traditional OS primitives.

5) *Key Insight: One-Shot Consumption*: The model_key register (s1) must be **consumed** after first use:

```
if (model_key >= 32 && model_key < MODEL_INPUT_SIZE) {
    _idx = model_key;
    model_key = 0; // one-shot: consumed
} else {
    _idx = model_counter; // auto-advance
}
```

Sub-model	Architecture	Role
Router	2->4->2->2	Task switching
Counter	255->256->256->255	Auto-incrementing state
Chargen	255->256->256->400	Character rendering
Game State	56->64->64->52	Squash physics
Renderer	48->128->256->300	Game rendering

Region	Offset	Size	Content
Router gates	0	2	gate_chargen, gate_squash
Chargers FB	4	400	20E20 character pixels
Squash state	404	52	ball_x, ball_y, paddle_y, etc.
Squash FB	456	300	20E15 game pixels
Counter state	768	255	counter argmax output

Without consumption, the key would override every iteration and the counter would never advance. This is analogous to interrupt flag clearing in traditional OS design.

G. AI-as-OS Taxonomy

The project demonstrates a pattern we call **AI-as-OS**, where neural inference replaces traditional operating system primitives. We taxonomy this approach along three dimensions:

- 1) *Execution Model*:
- 2) *Abstraction Layer*:
- 3) *Runtime Semantics*: **Key Insight**: The 141-line `runtime.c` replaces the entire OS. All application logic---counting, character rendering, game physics---is learned rather than programmed. This demonstrates that neural inference can serve as the primary computation model on minimal RISC-V cores.

H. Evaluation

1) *Benchmark Setup*: All benchmarks use the character generation workload (input character code to framebuffer glyph output). Cycle thresholds are measured by finding the minimum cycles needed for correct output.

2) *Optimization Results*: **Verilator backend (RTL): C++ backend (software)**:

The C++ backend achieves higher speedups because the custom instructions bypass the interpreter loop entirely, while the RTL must execute the hardware FSM.

Variant	Threshold Cycles	Speedup
<i>Verilator Backend (RTL)</i>		
Baseline (1-lane)	4,000,000	1.0x
x7b-8lane	400,000	10.0x
x7b-8lane-pmac	200,000	20.0x
x7b-8lane-pmac4	150,000	26.7x
<i>C++ Backend (Software)</i>		
Baseline	3,000,000	1.0x
x77 (optimized)	2,000	1,500x
x7b-8lane-pmac4	1,500	2,000x

Fig. 4. Benchmark Results: Speedup Comparison. The C++ backend achieves higher speedups (2000x) because custom instructions bypass the interpreter loop, while RTL must execute the hardware FSM (26.7x).

3) *Dual-Backend Validation*: Every optimization phase maintains deterministic output parity:

- Framebuffer bytes are identical across backends
- Counter state progresses identically

Key	Action	Neural Mechanism
Tab	Switch to squash mode	Router receives [0,1] -> gate [-]
Escape	Switch to chargen, squash pauses	Router receives [1,0], squash d
Backslash w/s	Switch to chargen, squash keeps running	Router receives [1,0], squash d
	Paddle up/down (squash mode)	Key injected via memory-mapp

Traditional OS	AI-as-OS (This Work)
Process scheduling	Router MLP switches between sub-models
Context switching	Model key consumption (one-shot)
Interrupt handling	Keyboard input neural tensor mapping
Memory management	Fixed-weight block-diagonal composition

- **Deterministic**: same input + same cycles = same output
This is enforced by automated CI tests that run both backends and compare outputs.

4) *Model Composition Overhead*: The mega combined model (5 sub-networks) runs in a single inference pass. Block-diagonal structure ensures zero cross-talk between sub-models. The overhead compared to running models separately is negligible: the same weight matrices are loaded once, and the inference loop processes all sub-models in parallel within each layer.

I. Conclusion

This project demonstrates that neural inference can serve as the primary execution model on a minimal RISC-V core. The key findings are:

1. **Descriptor-based neural ISA** provides a clean abstraction for MLP inference, separating firmware setup from hardware execution.
2. **Block-diagonal model composition** enables combining independently trained models without retraining, via a declarative glue JSON format.
3. **OS-like task switching** is achievable through a learned router MLP, demonstrating process scheduling semantics without traditional OS primitives.
4. **Dual-backend validation** ensures correctness across software and hardware implementations through automated differential tests.
5. **Significant speedups** are achievable: up to 2000x over software baselines (C++ backend) and 26.7x over unoptimized RTL (Verilator backend).

The repository contains a reproducible methodology: training pipeline, custom assembler, model compiler, differential verification, and phase-based benchmarks. This foundation enables further research into neural-first computing architectures.

J. Future Work

1) *Interrupt-Driven Neural Scheduling*: The current execution model is synchronous and loop-driven. Integrating timer and input IRQs would enable event-triggered neural inference, enabling evaluation of bounded latency and preemption behavior.

2) *Dual-Emulator Communication*: Running two emulator instances in lockstep with packet exchange would demonstrate neural multiplayer behavior, introducing synchronization and reproducibility questions for distributed deterministic simulation.

Traditional OS	AI-as-OS (This Work)
System calls	Neural instruction invocation
Device drivers	Input/output macro mapping
File system	N/A (weights are the "code")
Process table	Model map (sub-model registry)

Traditional OS	AI-as-OS (This Work)
Preemptive multitasking	Cooperative model switching
Shared memory	Separate output regions (block-diagonal)
Synchronization	Deterministic single-threaded execution
Resource allocation	Fixed-weight memory layout

Variant	Threshold Cycles	Speedup vs Baseline
Baseline (1-lane)	4,000,000	1.0x
x7b-8lane	400,000	10.0x
x7b-8lane-pmac	200,000	20.0x
x7b-8lane-pmac4	150,000	26.7x

Variant	Threshold Cycles	Speedup vs Baseline
Baseline	3,000,000	1.0x
x77 (optimized)	2,000	1,500x
x7b-8lane-pmac4	1,500	2,000x

3) *Hardware I/O Abstraction*: Char-device streams and MMIO/register-mapped peripherals with small software drivers would bridge input events to neural tensors and neural outputs to actuator/framebuffer paths.

K. Acronyms

- ABI: Application Binary Interface
- CI: Continuous Integration
- CTest: CMake test driver
- ELF: Executable and Linkable Format
- FC: Fully Connected (neural layer)
- ISA: Instruction Set Architecture
- MAC: Multiply-Accumulate
- MLP: Multi-Layer Perceptron
- PMAC: Parallel Multiply-Accumulate
- RTL: Register Transfer Level

L. References

- [1] RISC-V International, "The RISC-V Instruction Set Manual," <https://riscv.org/technical/specifications/>
- [2] Verilator project, "Verilator Open-Source SystemVerilog Simulator," <https://www.veripool.org/verilator/>
- [3] PyTorch contributors, "PyTorch: An Imperative Style, High-Performance Deep Learning Library," NeurIPS, 2019.
- [4] GNU Binutils, "GNU assembler (as) and objcopy documentation," <https://sourceware.org/binutils/>
- [5] Repository artifact, `documentation/collected-data/phase25-neural-lane-cycle-comparison.json`.
- [6] Repository source, `projects/rv32_assembler/README.md` and `projects/rv32_assembler/src/{parser.rs,encoder.rs}`.
- [7] Repository source, `projects/emulator/COMBINING.md`, `projects/emulator/combined-mlp.md`.
- [8] S. Schwartz-Ziv et al., "FS-Merge: Foldable SuperNets for Efficient Model Merging," arXiv 2024.
- [9] T. Wang et al., "LoRA-LEGO: Concatenation-Summation Equivalence for LoRA Merging," arXiv 2024.
- [10] Arcee AI, "mergekit: A Toolkit for Merging Large Language Models," GitHub, 2023.
- [11] R. Price, "nCPU: A Fully Differentiable Neural CPU," 2025.
- [12] "RV-SCNN: Custom Neural ISA Extensions for RISC-V CNN/SNN Inference," IEEE TCAD, 2025.
- [13] "MARVEL: Model-Class-Specific RISC-V Extensions via ASIP Designer," arXiv, 2025.
- [14] "FPGA-Accelerated RISC-V ISA Extensions for Efficient Neural Network Inference on Edge Devices," arXiv, 2025.
- [15] "Vmxdotp: A RISC-V Vector ISA Extension for Efficient Microscaling (MX) Format Acceleration," arXiv, 2026.
- [16] "Mixed-precision Neural Networks on RISC-V Cores: ISA extensions for Multi-Pumped Soft SIMD Operations," arXiv, 2024.
- [17] "Optimised Extension of an Ultra-Low-Power RISC-V Processor to Support Lightweight Neural Network Models," MDPI, 2025.